

# Artificial Intelligence ——总结



Yanghui Rao

Assistant Prof., Ph.D

School of Data and Computer Science,

Sun Yat-sen University

raoyangh@mail.sysu.edu.cn

# Probability

- **Product rule:**

$$P(A, B) = P(A)P(B | A) = P(B)P(A | B)$$

$$P(A, B_1, B_2, B_3) = P(A)P(B_1 | A)P(B_2 | A, B_1)P(B_3 | A, B_1, B_2)$$

- **Sum rule:**  $P(A) = P(A, B) + P(A, B^c)$

$$P(A) = \sum_{i=1}^n P(A, B_i)$$

$$= \sum_{i=1}^n P(A | B_i)P(B_i)$$

# Probability

- Exercise: 假设有一盒骰子，里面有4面的（点数为1、2、3、4），6面的、8面的、12面的、20面的均匀骰子各1个。如果我随机从盒子中选一个骰子，投掷它得到了点数5。那么我选中的骰子为4面、6面、8面、12面、20面的概率各是多少？

**答案:** 0, 0.392, 0.294, 0.196, 0.118

# Probability

- There are two random variables  $X$  and  $Y$ . Which of the following is always true?
  - A.  $\sum_X P(X|Y) = 1$
  - B.  $\sum_Y P(X|Y) = 1$
  - C. All of the above
  - D. None of the above
- Is the statement True or False? Entropy of a discrete random variable is always non-negative.

# Probability

- There are two random variables X and Y. Which of the following is always true? (Answer: A)
  - A.  $\sum_X P(X|Y) = 1$
  - B.  $\sum_Y P(X|Y) = 1$
  - C. All of the above
  - D. None of the above
- Is the statement True or False? Entropy of a discrete random variable is always non-negative. (Answer: True)

# Truth Tables

- Truth tables are used to define logical connectives and to determine when a complex sentence is true given the values of the symbols in it
- Note that  $\Rightarrow$  is a logical connective, so  $P \Rightarrow Q$  is a logical sentence and has a truth value, i.e., is either true or false

*Truth tables for the five logical connectives*

<b>P</b>	<b>Q</b>	<b><math>\neg P</math></b>	<b><math>P \wedge Q</math></b>	<b><math>P \vee Q</math></b>	<b><math>P \Rightarrow Q</math></b>	<b><math>P \Leftrightarrow Q</math></b>
False	False	True	False	False	True	True
False	True	True	False	True	True	False
True	False	False	False	True	False	False
True	True	False	True	True	True	True

# Quantifier Scope

- If a quantifier  $Q$  is followed by  $($ , then the scope of  $Q$  is to the matched  $)$ 
  - $\forall x (F(x) \Leftrightarrow F(h))$
- If a quantifier  $Q$  is not followed by  $($  or another quantifier, then the scope of  $Q$  is to the first connective
  - $\forall x F(x) \Leftrightarrow F(h)$
- If a quantifier  $Q1$  is followed by another quantifier  $Q2$ , then the scope of  $Q1$  is to the scope of  $Q2$ 
  - $\forall x \exists y R(x, y)$
- $F$ : ... can fly
- $h$ : human being

False		True
$\forall x (F(x) \Leftrightarrow F(h))$	$\nLeftrightarrow$	$\forall x F(x) \Leftrightarrow F(h)$

# Exercise

- Fill in the following truth table:

P	Q	$(P \Rightarrow Q) \wedge (Q \Rightarrow P)$	$(\neg P \vee Q) \Leftrightarrow (P \Rightarrow Q)$
True	True		
True	False		
False	True		
False	False		

- If we represent “... is hot” by  $H(\dots)$ , and represent “fire” by  $f$ , what are the values of “ $H(f) \Rightarrow \forall x H(x)$ ” and “ $\exists x (H(f) \Leftrightarrow H(x))$ ”?



# Exercise


- Fill in the following truth table:

P	Q	$(P \Rightarrow Q) \wedge (Q \Rightarrow P)$	$(\neg P \vee Q) \Leftrightarrow (P \Rightarrow Q)$
True	True	True	True
True	False	False	True
False	True	False	True
False	False	True	True

- If we represent “... is hot” by  $H(\dots)$ , and represent “fire” by  $f$ , what are the values of “ $H(f) \Rightarrow \forall x H(x)$ ” and “ $\exists x (H(f) \Leftrightarrow H(x))$ ”? (Answer: False, True)

# $k$ -Nearest Neighbor

$k$ -NN using a majority voting scheme



# Naïve Bayesian Classifier

- This can be derived from Bayes' theorem

$$P(C_i | \mathbf{X}) = \frac{P(\mathbf{X} | C_i)P(C_i)}{P(\mathbf{X})}$$

- Since  $P(\mathbf{X})$  is constant for all classes, only

$$P(C_i | \mathbf{X}) \propto P(\mathbf{X} | C_i)P(C_i)$$

needs to be maximized

- $P(C_i)$  can be obtained from training set  $s_i/s$

# Derivation

- **Assumption:** attributes are conditionally independent (i.e., no dependence relation between attributes):  
$$P(\mathbf{X} | C_i) = \prod_{k=1}^n P(x_k | C_i)$$
- This greatly reduces the computation cost:  
Only counts the class distribution
- If  $A_k$  is categorical,  $P(x_k | C_i) = s_{ik}/s_i$ , count the distribution
- If  $A_k$  is continuous-valued,  $P(x_k | C_i)$  can be computed based on Gaussian distribution

# Exercise

- What is the meaning of “k” for the k-Nearest Neighbor (i.e., k-NN) and the k-Means clustering algorithm?
- If using k-NN for classification, what is the predicted class label when “ $x = 5$ ”? Is there any difference if based on City Block, Euclidean, or Supremum distance?

Note: Given a testing sample  $x$ , if there are multiple training samples' distances are the nearest, k-NN classifier will use the mode (众数) of the class labels of all nearest training samples as the predicted class label of  $x$

X	Y
2	-
3.3	-
3.2	-
3.1	+
5	?

X	Y
2	+
3.3	-
3.2	-
3.1	+
5	?

# Exercise

- What is the meaning of “k” for the k-Nearest Neighbor (i.e., k-NN) and the k-Means clustering algorithm?
  - Answer: a) The parameter “k” means the number of neighbors used to classify test examples for the k-NN. b) The parameter “k” specifies the number of clusters for the k-Means.
- Given a testing sample  $x$ , if there are multiple training samples' distances are the nearest, k-NN classifier will use the mode (众数) of the class labels of all nearest training samples as the predicted class label of  $x$ .
  - Answer: There is no difference if based on those distance measures. (1) Left table. The predicted class label is “-”; (2) Right table. The predicted class label is “-” for  $k=1, 2, 3$  and “Unknown” for  $k > 3$ .

# Information Gain (ID3)

- Class label: buy\_computer="yes/no"
- 用字母 $D$ 表示类标签，字母 $A$ 表示每个属性
- $H(D)=0.940$  14个训练样本中，9个买了电脑

$$H(D) = -\frac{9}{14} \log_2 \frac{9}{14} - \left(1 - \frac{9}{14}\right) \log_2 \left(1 - \frac{9}{14}\right)$$

- $H(D | A = "age") = 0.694$

$$\begin{aligned} H(D | A = "age") &= \frac{5}{14} \times \left( -\frac{2}{5} \log_2 \frac{2}{5} - \frac{3}{5} \log_2 \frac{3}{5} \right) \\ &+ \frac{4}{14} \times \left( -\frac{4}{4} \log_2 \frac{4}{4} - \frac{0}{4} \log_2 \frac{0}{4} \right) + \frac{5}{14} \times \left( -\frac{3}{5} \log_2 \frac{3}{5} - \frac{2}{5} \log_2 \frac{2}{5} \right) \end{aligned}$$

# Information Gain (ID3)

- Class label: buy\_computer="yes/no"
  - Compute the mutual information (互信息) between  $D$  and each attribute  $A$
  - $H(D)=0.940$
  - $H(D|A="age")=0.694$
  - $g(D,A="age")=0.246$
  - $g(D,A="income")=0.029$
  - $g(D,A="student")=0.151$
  - $g(D,A="credit\_rating")=0.048$
- “age”这个属性的条件熵最小（等价于信息增益最大），因而首先被选出作为根节点**
- |  |              |
|--|--------------|
|  | $g(D, A)$    |
|  | $= H(D)$     |
|  | $- H(D   A)$ |



# Information Gain Ratio (C4.5)

- $\text{GainRatio}_A(D) = \text{Gain}_A(D) / \text{SplitInfo}_A(D)$

$$\text{SplitInfo}_A(D) = - \sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2 \left( \frac{|D_j|}{|D|} \right)$$

- $\text{GainRatio}_{A=\text{"income"}}(D) = ?$

$$\text{SplitInfo}_{A=\text{"income"}}(D)$$

$$\begin{aligned} &= -\frac{4}{14} \times \log_2 \left( \frac{4}{14} \right) - \frac{6}{14} \times \log_2 \left( \frac{6}{14} \right) - \frac{4}{14} \times \log_2 \left( \frac{4}{14} \right) \\ &= 0.926 \end{aligned}$$

- $\text{GainRatio}_{A=\text{"income"}}(D) = 0.029 / 0.926 = 0.031$

# Gini Index (CART)

- $D$  has 9 samples in `buys_computer` = “yes” and 5 in “no”

$$gini(D) = 1 - \left(\frac{9}{14}\right)^2 - \left(\frac{5}{14}\right)^2 = 0.459$$

- The attribute *income* partitions  $D$  into 10 in  $D_1$ : {medium, high} and 4 in  $D_2$

$$gini_{income \in \{\text{medium, high}\}}(D) = \frac{10}{14} gini(D_1) + \frac{4}{14} gini(D_2)$$

$$= \frac{10}{14} \left( 1 - \left(\frac{6}{10}\right)^2 - \left(\frac{4}{10}\right)^2 \right) + \frac{4}{14} \left( 1 - \left(\frac{1}{4}\right)^2 - \left(\frac{3}{4}\right)^2 \right)$$

$$= 0.450 = gini_{income \in \{\text{low}\}}(D)$$

# Decision Tree

- But how can we compute the gini index, information gain of an attribute that is **continuous-valued**?
  - Given  $v$  values of  $A$ , then  $v-1$  possible splits are evaluated. For example, the midpoint between the values  $a_i$  and  $a_{i+1}$  of  $A$  is  $(a_i + a_{i+1}) / 2$

# Incorporating model complexity

- In the case of a decision tree, let
  - $L$  be the number of leaf nodes.
  - $n_l$  be the  $l$ -th leaf node.
  - $m(n_l)$  be the number of training records classified by  $n_l$ .
  - $r(n_l)$  be the number of misclassified records by  $n_l$ .
  - $\zeta(n_l)$  be a penalty term associated with the node  $n_l$ .
- The resulting error  $e_c$  of the decision tree can be estimated as follows:

$$e_c = \frac{\sum_{l=1}^L (r(n_l) + \zeta(n_l))}{\sum_{l=1}^L m(n_l)}$$

# Exercise

- We consider the training examples shown in the following table for a binary classification problem.
  - Calculate the respective changes in the Gini index value when  $a_1$  and  $a_2$  are used for partitioning the training set.
  - Calculate the respective changes in the classification (**training**) error when  $a_1$  and  $a_2$  are used for partitioning the training set.

$a_1$	$a_2$	$a_3$	Target Class
T	T	1	+
T	T	6	+
T	F	5	-
F	F	4	+
F	T	7	-
F	T	3	-
F	F	8	-
T	F	7	+
F	T	5	-

# Exercise

- (1) The original Gini index is  $1 - (\frac{4}{9})^2 - (\frac{5}{9})^2 = 0.494$

After splitting on  $a_1$ , the Gini index becomes

$$\frac{4}{9}[1 - (\frac{3}{4})^2 - (\frac{1}{4})^2] + \frac{5}{9}[1 - (\frac{1}{5})^2 - (\frac{4}{5})^2] = 0.344$$

As a result, the change in Gini index is

$$\Delta G(a_1) = 0.494 - 0.344 = 0.15.$$

After splitting on  $a_2$ , the Gini index becomes

$$\frac{5}{9}[1 - (\frac{2}{5})^2 - (\frac{3}{5})^2] + \frac{4}{9}[1 - (\frac{2}{4})^2 - (\frac{2}{4})^2] = 0.489$$

As a result,

$$\Delta G(a_2) = 0.494 - 0.489 = 0.005.$$

# Exercise

- (2) The original classification error is  $1 - \max(\frac{4}{9}, \frac{5}{9}) = \frac{4}{9}$

After splitting on  $a_1$ , the classification error becomes

$$\frac{4}{9}[1 - \max(\frac{3}{4}, \frac{1}{4})] + \frac{5}{9}[1 - \max(\frac{1}{5}, \frac{4}{5})] = \frac{2}{9}$$

As a result, the change in classification error is

$$\Delta E(a_1) = 4/9 - 2/9 = 2/9.$$

After splitting on  $a_2$ , the classification error becomes


$$\frac{5}{9}[1 - \max(\frac{2}{5}, \frac{3}{5})] + \frac{4}{9}[1 - \max(\frac{2}{4}, \frac{2}{4})] = \frac{4}{9}$$

As a result,

$$\Delta E(a_2) = 4/9 - 4/9 = 0.$$

# Exercise

- Consider the following decision tree with four nodes A, B, C, D and five leaf nodes:




- What is the original Gini index value of the data set?
- What is the value of the penalty term for each leaf node if the generalization error is 0.5?
- Suppose a penalty term of 1 is assigned to each leaf node, estimate the generalization error if the sub-tree associated with node C is pruned and replaced with a leaf node.



# Exercise

- Consider the following decision tree with four nodes A, B, C, D and five leaf nodes:



- What is the original Gini index value of the data set? **0.5**

$$1 - \left(\frac{1}{2}\right)^2 - \left(\frac{1}{2}\right)^2 = 0.5$$

- What is the value of the penalty term for each leaf node if the generalization error is 0.5? **1**

$$\frac{10 + 1 \times 5}{30} = 0.5$$

- Suppose a penalty term of 1 is assigned to each leaf node, estimate the generalization error if the sub-tree associated with node C is pruned and replaced with a leaf node. **0.5**

$$\frac{12 + 1 \times 3}{30} = 0.5$$

# Linear Regression

- Gradient descent solution?

$$n^{-1} \sum_{i=1}^n (y_i - w_0 - w_1 x_i) = 0$$

$$n^{-1} \sum_{i=1}^n x_i (y_i - w_0 - w_1 x_i) = 0$$

$$Q(w_0, w_1) = \min_{w_0, w_1} \sum_{i=1}^n (y_i - w_0 - w_1 x_i)^2$$

$$\partial Q(w_0, w_1) / \partial w_0 = 0$$


$$\partial Q(w_0, w_1) / \partial w_1 = 0$$

$$-2 \sum_{i=1}^n (y_i - w_0 - w_1 x_i) = 0$$

$$-2 \sum_{i=1}^n x_i (y_i - w_0 - w_1 x_i) = 0$$

# Perceptron Learning Algorithm

- Difficult: the set of  $h(\mathbf{x})$  is of infinite size
- Idea: start from some initial weight vector  $\mathbf{w}_{(0)}$ , and “correct” its mistakes on  $D$
- For  $t = 0, 1, \dots$ 
  - find a mistake of  $\mathbf{w}_{(t)}$  called  $(\mathbf{x}_{n(t)}, y_{n(t)})$   
 $\text{sign}(\mathbf{w}_{(t)}^T \mathbf{x}_{n(t)}) \neq y_{n(t)}$
  - (try to) correct the mistake by  
$$\mathbf{w}_{(t+1)} \leftarrow \mathbf{w}_{(t)} + y_{n(t)} \mathbf{x}_{n(t)}$$
  - until no more mistakes
- Return last  $\mathbf{W}$  (called  $\mathbf{W}_{\text{PLA}}$ )



# Perceptron Learning Algorithm

- Only if there exists an hyperplane that correctly classifies the data, the Perceptron procedure is guaranteed to converge; furthermore, the algorithm may give different results depending on the order in which the elements are processed, indeed several different solutions exist.

# Exercise


- What are the values of weights  $w_0$ ,  $w_1$ , and  $w_2$  for the perceptron whose decision surface is illustrated in the Figure? Assume the surface crosses the  $x_1$  axis at -1, and the  $x_2$  axis at 2.

- Answer:

$w_0 =$

$w_1 =$

$w_2 =$



# Exercise

- What are the values of weights  $w_0$ ,  $w_1$ , and  $w_2$  for the perceptron whose decision surface is illustrated in the Figure? Assume the surface crosses the  $x_1$  axis at -1, and the  $x_2$  axis at 2.

The surface crosses  $(-1, 0)$  and  $(0, 2)$

One surface:  $-1 - x_1 + 0.5 \cdot x_2 = 0$


- Answer:

$$w_0 = -1 \cdot C$$

$$w_1 = -1 \cdot C$$

$$w_2 = 0.5 \cdot C$$


(where  $C > 0$ )



# Logistic Regression Model

- The logistic distribution constrains the estimated probabilities to lie between 0 and 1.
- The estimated probability  $p(y=1 | \mathbf{X})$  is:

$$p = \frac{1}{1 + e^{-w_0 - \sum_{j=1}^d w_j x_j}} = \frac{e^{w_0 + \sum_{j=1}^d w_j x_j}}{1 + e^{w_0 + \sum_{j=1}^d w_j x_j}}$$
$$= \frac{1}{1 + e^{-\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}}} = \frac{e^{\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}}}{1 + e^{\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}}}$$



- if you let  $w_0 + \sum_{j=1}^d w_j x_j = 0$ , then  $p = 0.5$
- as  $w_0 + \sum_{j=1}^d w_j x_j$  gets really big,  $p$  approaches 1
- as  $w_0 + \sum_{j=1}^d w_j x_j$  gets really small,  $p$  approaches 0

**PLA ?**

# Logistic Regression Model

- The likelihood function is  $\prod_{i=1}^n (p_i)^{y_i} (1 - p_i)^{1-y_i}$
- We want to maximize the log likelihood:

$$L(\tilde{\mathbf{W}}) = \sum_{i=1}^n (y_i \log p_i + (1 - y_i) \log(1 - p_i))$$

$$= \sum_{i=1}^n \left( y_i \log \frac{p_i}{1 - p_i} + \log(1 - p_i) \right)$$

$$= \sum_{i=1}^n \left( y_i \tilde{\mathbf{W}}^T \tilde{\mathbf{X}}_i - \log(1 + e^{\tilde{\mathbf{W}}^T \tilde{\mathbf{X}}_i}) \right)$$

$$\frac{\partial L(\tilde{\mathbf{W}})}{\partial \tilde{\mathbf{W}}} = \sum_{i=1}^n \left[ \left( y_i - \frac{e^{\tilde{\mathbf{W}}^T \tilde{\mathbf{X}}_i}}{1 + e^{\tilde{\mathbf{W}}^T \tilde{\mathbf{X}}_i}} \right) \tilde{\mathbf{X}}_i \right]$$

- It is equal to minimize the cost function

$$C(\tilde{\mathbf{W}}) = -L(\tilde{\mathbf{W}}) = -\sum_{i=1}^n (y_i \log p_i + (1 - y_i) \log(1 - p_i)) \quad \text{Cross-entropy}$$



# Logistic Regression Model

- Gradient Decent (梯度下降)

- Calculate the gradient vector
- Update the weighting in the opposite direction of the gradient vector at each surface point

- Repeat:  $\tilde{\mathbf{W}}_{new}^{(j)} = \tilde{\mathbf{W}}^{(j)} - \eta \frac{\partial C(\tilde{\mathbf{W}})}{\partial \tilde{\mathbf{W}}^{(j)}}$ 
$$= \tilde{\mathbf{W}}^{(j)} - \eta \sum_{i=1}^n \left[ \left( \frac{e^{\tilde{\mathbf{W}}^T \tilde{\mathbf{X}}_i}}{1 + e^{\tilde{\mathbf{W}}^T \tilde{\mathbf{X}}_i}} - y_i \right) \tilde{\mathbf{X}}_i^{(j)} \right]$$
- Until convergence

# Neural Network

- Given a unit  $j$  in a hidden or output layer, the net input,  $I_j$ , to unit  $j$  is  $I_j = \sum_i w_{ij} O_i + \theta_j$

where  $w_{ij}$  is the weight of the connection from unit  $i$  in the previous layer to unit  $j$ ;  $O_i$  is the output of unit  $i$  from the previous layer; and  $\theta_j$  is the bias of the unit.

- Given the net input  $I_j$  to unit  $j$ , then  $O_j$ , the output of unit  $j$ , is computed as  $O_j = \frac{1}{1 + e^{-I_j}}$

- For a unit  $k$  in the output layer, the error  $Err_k$  is computed by

$$Err_k = O_k (1 - O_k) (T_k - O_k)$$


- The error of a hidden layer unit  $j$  is

$$Err_j = O_j (1 - O_j) \sum_k Err_k w_{jk}$$

- Weights are updated by

$$w_{jk} = w_{jk} + \eta Err_k O_j$$

$$\theta_k = \theta_k + \eta Err_k$$



Propagate the  
inputs forward


Backpropagate  
the error

# Neural Network

- Minimize the error of node  $O_k$
- We define it as  $E = \frac{1}{2}e^2 = \frac{1}{2}(T - O)^2$
- To adjust weight  $w_{jk}$ , we first calculate the partial derivation of  $E$  on  $w_{jk}$

$$\begin{aligned}\frac{\partial E}{\partial w_{jk}} &= \frac{\partial E}{\partial e} \times \frac{\partial e}{\partial O_k} \times \frac{\partial O_k}{\partial w_{jk}} \\ &= -(e) \times (O_k(1 - O_k)) \times (O_j) \\ &= -(T_k - O_k)O_k(1 - O_k)O_j\end{aligned}$$

- and then use the “gradient decent”



# Apriori Algorithm

- **自连接**: 用  $L_{k-1}$  自连接得到  $C_k$
- **修剪**: 一个  $k$ -项集, 如果他的一个  $k-1$  项集 (他的子集) 不是频繁的, 那他本身也不可能是频繁的。
- pseudo code:

$C_k$ : Candidate itemset of size  $k$

$L_k$ : frequent itemset of size  $k$

$L_1 = \{\text{frequent items}\};$

**for** ( $k = 1; L_k \neq \emptyset; k++$ ) **do begin**

$C_{k+1}$  = candidates generated from  $L_k$ ;

**for each** transaction  $t$  in database **do**

        increment the count of all candidates in  $C_{k+1}$  that are contained in  $t$

$L_{k+1}$  = candidates in  $C_{k+1}$  with *minsup*


**end**

**return**  $\cup_k L_k$ ;

# Maximal Frequent Itemsets

- A maximal frequent itemset is defined as a frequent itemset for which **none of its immediate supersets are frequent.**

# Maximal Frequent Itemsets



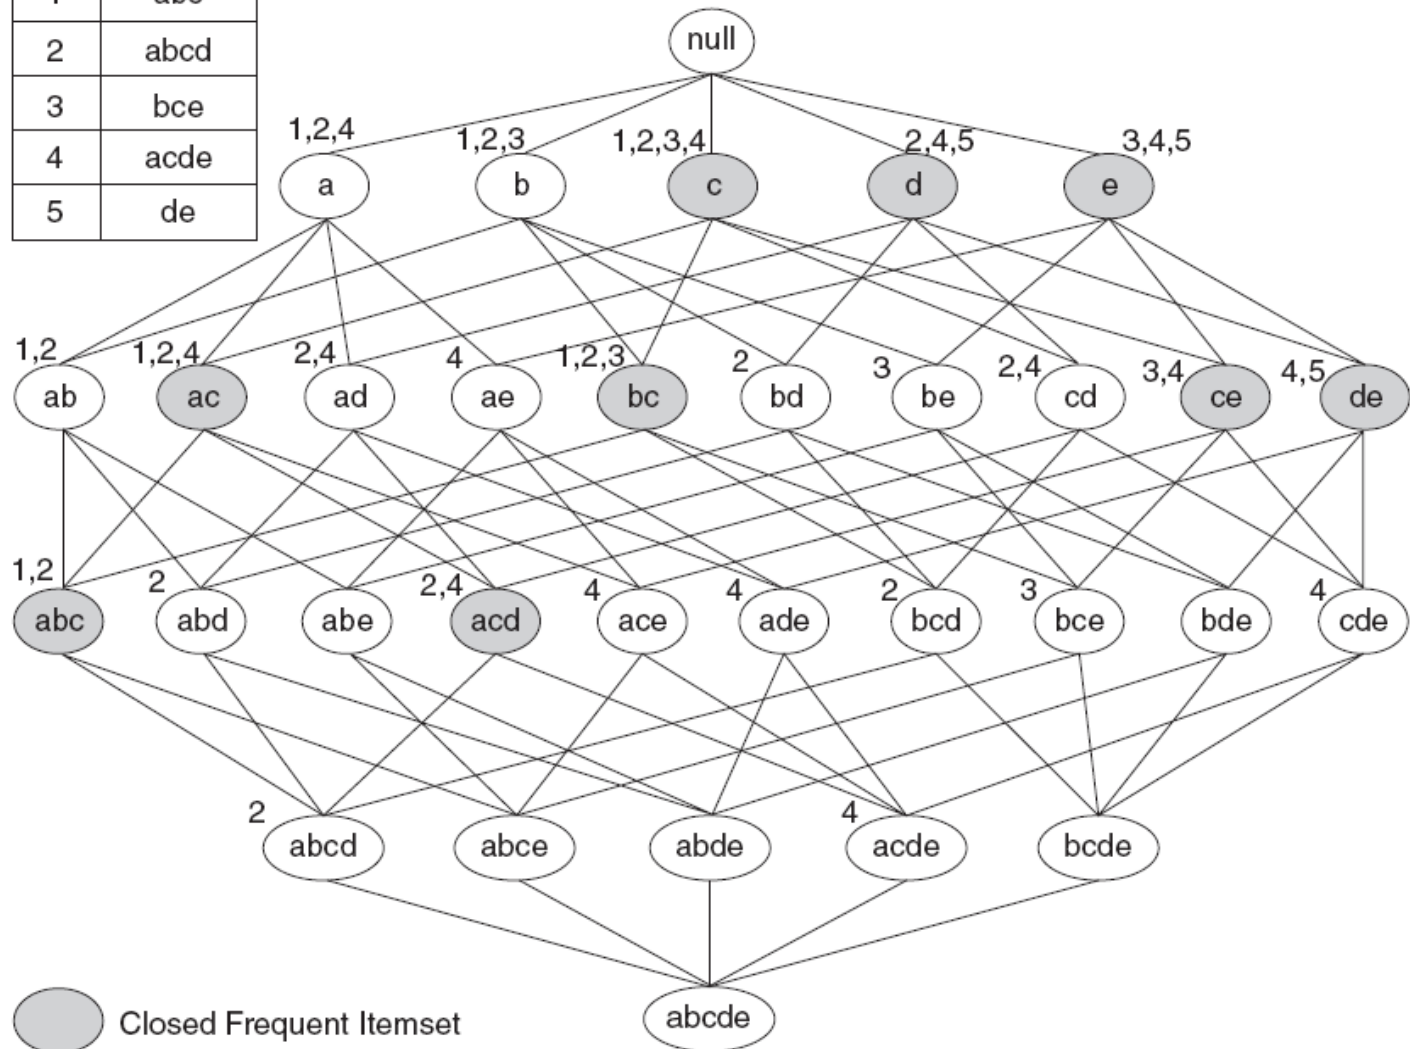
# Closed Frequent Itemsets

- An itemset  $X$  is closed if none of its immediate supersets has exactly the same support count as  $X$ .
- In other words,  $X$  is not closed if at least one of its immediate supersets has the same support count as  $X$ .

# Closed Frequent Itemsets

TID	Items
1	abc
2	abcd
3	bce
4	acde
5	de

minsup = 40%






# Partitional Clustering

- $k$ -Means: Repeat...
  - Choose  $k$  arbitrary '**centroids**'
  - Assign each document to nearest centroid
  - Re-compute centroids
- **Example of  $k$ -Means (划分法)**

# DBSCAN

- We need to classify a point as being
  - In the interior of a dense region (a **core** point, 核心点).
  - At the edge of a dense region (a **border** point, 边界点)
  - In a sparsely occupied region (a **noise** or background point, 噪音点).
- The concepts of core, border and noise points are illustrated as follows.


# DBSCAN




# Hierarchical Clustering

- A hierarchical clustering is often displayed graphically using a tree-like diagram called the dendrogram (树状图).
- The dendrogram displays both
  - the cluster-subcluster relationships and
  - the order in which the clusters are merged (agglomerative) or split (divisive).
- For sets of 2-D points, a hierarchical clustering can also be graphically represented using a nested cluster diagram.

# Hierarchical Clustering



(a) Dendrogram.



(b) Nested cluster diagram.

# Hierarchical Clustering

- Different definitions of cluster distance leads to different versions of hierarchical clustering.
- These versions include
  - Single link (单连接) or MIN
  - Complete link (全连接) or MAX
  - Group average (组平均)

# Single Link

- We now consider the single link or MIN version of hierarchical clustering.
- In this case, the distance of two clusters is defined as the minimum of the distance between any two points in the two different clusters.
- This technique is good at handling non-elliptical (非球状的) shapes.

# Complete Link

- We now consider the complete link or MAX version of hierarchical clustering.
- In this case, the distance of two clusters is defined as the maximum of the distance between any two points in the two different clusters.
- Complete link tends to produce clusters with globular (球状) shapes.